

LOW-RANK APPROACHES FOR CONTINUAL LEARNING IN LLMS

THESIS

Submitted in Partial Fulfillment of
the Requirements for
the Degree of

MASTER OF SCIENCE (Computer Engineering)

at the

**NEW YORK UNIVERSITY
TANDON SCHOOL OF ENGINEERING**

by

Harshal Kulkarni

May 2025

LOW-RANK APPROACHES FOR CONTINUAL LEARNING IN LLMS

THESIS

Submitted in Partial Fulfillment of

the Requirements for

the Degree of

MASTER OF SCIENCE (Computer Engineering)

at the

NEW YORK UNIVERSITY

TANDON SCHOOL OF ENGINEERING

by

Harshal Kulkarni

May 2025

Approved:




Department Chair Signature

May 7, 2025

Date

Approved by the Guidance Committee:

Major: **Computer Engineering**



Prof. Anna Choromanska

Associate Professor

New York University, Tandon School of
Engineering

05/07/2025

Date

Microfilm or other copies of this thesis are obtainable from:

UMI Dissertation Publishing

ProQuest CSA

789 E. Eisenhower Parkway

P.O. Box 1346

Ann Arbor, MI 48106-1346

Table of Contents

Table of Contents	iii
1 Introduction	1
2 Background	4
2.1 Continual Learning and Catastrophic Forgetting	4
2.2 Continual Learning in LLMs	6
2.3 Parameter-Efficient Fine-Tuning (PEFT)	6
2.4 Continual Learning with LoRA	7
3 Orthogonal Low-Rank Adaptation	9
3.1 Motivation and Challenges	9
3.2 Orthogonal LoRA (O-LoRA)	10
3.3 O-LoRA Architecture	11
3.4 Empirical Results	11
3.5 Modified Orthogonal LoRA (M-LoRA)	13
3.6 Empirical Results	14
4 Task Arithmetic for Model Merging	17
4.1 The Concept of Task Vectors	17
4.2 Theoretical Insights on Task Disentanglement	18
4.3 Orthogonality of Task Vectors	19
4.4 Orthogonality Regularization for Task Arithmetic	20

4.5	<i>Empirical Results</i>	20
5	Experimental Setup and Results	23
5.1	<i>Model and Framework</i>	23
5.2	<i>Datasets and Tasks</i>	24
5.3	<i>Evaluation Metrics</i>	24
5.4	<i>Method Comparison</i>	25
5.5	<i>Performance Summary</i>	25
6	Discussion	27
6.1	<i>Stability vs. Plasticity Trade-Off</i>	27
6.2	<i>Orthogonality in Weight Space</i>	28
6.3	<i>Effectiveness of Task Arithmetic</i>	28
6.4	<i>Comparative Observations</i>	29
6.5	<i>Guidelines for Method Selection</i>	29
6.6	<i>Limitations and Future Work</i>	30
7	Conclusion and Future Work	31
7.1	<i>Conclusion</i>	31
7.2	<i>Future Work</i>	32
	Bibliography	33

VITA

Harshal Kulkarni was born in Pune, India. He received his Bachelor of Engineering degree in Electronics and Communication from BITS Pilani, India.

From January 2021 to July 2023, he was employed at Samsung R&D Institute, Bangalore, where he contributed to the development of backend systems.

He joined New York University Tandon School of Engineering in Fall 2023 to pursue a Master of Science in Computer Engineering with a specialization in Machine Learning. His graduate studies emphasized Machine Learning, deep learning, continual learning, and large language models (LLMs).

From Fall 2023 to Spring 2025, he served as a Graduate Teaching Assistant for the Machine Learning course, supporting instruction in regression, classification, SVMs, kernels, generative models, and HMMs. He conducted his thesis research on low-rank continual learning for LLMs under the supervision of Prof. Anna Choromanska, with collaboration and experimental support from Kristi Topollai and Tolga Dimlioglu.

This research was carried out in the Machine Learning and Systems Lab at NYU and was supported in part by academic research resources provided by the NYU High Performance Computing (HPC) facility.

ACKNOWLEDGEMENTS

I would like to express my sincere gratitude to my advisor, Prof. Anna Choromanska, for her guidance, support, and insightful feedback throughout this thesis.

Special thanks to Kristi Topollai and Tolga Dimlioglu for their collaboration and contributions to the research experiments.

I also extend my appreciation to the faculty and peers at NYU Tandon for creating a stimulating and supportive research environment.

Finally, I thank my family for their constant encouragement and belief in my journey.

ABSTRACT

LOW-RANK APPROACHES FOR CONTINUAL LEARNING IN LLMS

by

Harshal Kulkarni

Advisor: Prof. Anna Choromanska

Submitted in Partial Fulfillment of the Requirements for
the Master of Science (Computer Engineering)

May 2025

Large Language Models (LLMs) excel in natural language tasks but suffer from catastrophic forgetting when trained continually on sequential tasks. This thesis introduces low-rank adaptation methods that are both parameter-efficient and orthogonally constrained to preserve task knowledge. We evaluate Orthogonal LoRA (O-LoRA) and its modified form on standard NLP benchmarks, demonstrating reduced forgetting and improved multitask generalization. Additionally, we explore task arithmetic and its theoretical underpinnings in neural weight space geometry. Our results show that orthogonality promotes disentangled task vectors and enables successful model merging. These insights push forward the practical and theoretical understanding of continual learning for LLMs.

List of Figures

1.1	Illustration of catastrophic forgetting in continual learning: performance on Task A drops sharply after training on Task B.	3
2.1	Continual learning strategies: (a–b) regularization with Elastic Weight Consolidation (EWC), and (c) parameter-isolated architectural method (Progressive Networks).	5
2.2	Six categories of continual learning in LLMs. These differ based on task visibility, data format, and supervision signals.	6
2.3	Common parameter-efficient fine-tuning (PEFT) methods including LoRA, DyLoRA, and DoRA. These methods reduce the number of trainable parameters per task.	7
3.1	Visualization of Orthogonal Gradient Descent (OGD): Gradients for new tasks are projected onto subspaces orthogonal to previously learned task gradients to reduce forgetting. O-LoRA adapts this principle using orthogonality between LoRA adapters.	11
3.2	O-LoRA architecture: Each task-specific LoRA update is constrained to be orthogonal to previous tasks’ update directions, reducing interference and enabling sequential learning.	12
3.3	M-LoRA extends O-LoRA by constraining the actual task-specific update matrices (AB) to lie in orthogonal subspaces.	14
3.4	Cosine similarity distribution among task vectors $A_i B_i$. The majority cluster around zero, indicating emergent orthogonality.	15

4.1	Task Arithmetic concept: Individual task-specific LoRA updates are merged to create a single model that performs well on all tasks.	18
4.2	Disentangled task vector behavior: Separate task-specific directions influence disjoint input regions, enabling effective merging via Task Arithmetic.	19
4.3	Comparison of cosine similarity heatmaps of task vectors with and without orthogonality regularization. Regularization (a) reduces overlap between task vectors, improving merge stability; without it (b), interference is more likely [11].	21

List of Tables

3.1	Performance of O-LoRA on benchmark datasets (OPT-1.3B)	12
3.2	Performance of sequential adaptation methods on benchmark datasets (OPT-1.3B)	14
4.1	Performance of Task Arithmetic (OPT-1.3B) with varying interpolation weight α	21
5.1	Final accuracy (%) per task and average across methods (OPT-1.3B)	26

Introduction

Large Language Models (LLMs) have revolutionized natural language processing (NLP), delivering state-of-the-art results on tasks such as text classification, summarization, question answering, and dialog generation. However, their performance degrades when they are fine-tuned sequentially on multiple tasks — a challenge known as *catastrophic forgetting* [7].

Retraining the full model from scratch for every task is memory-intensive and computationally expensive. Storing a separate fine-tuned model for each task is also impractical at the scale of modern LLMs. Thus, there is a pressing need for strategies that retain and integrate knowledge incrementally — without accessing prior data and while remaining efficient in memory and compute.

This thesis addresses the central problem of continual learning in LLMs — the ability to acquire new tasks sequentially without degrading performance on previously learned tasks.

Research Questions

We investigate the following key questions:

- Can task-specific low-rank updates be constrained to occupy disjoint subspaces to reduce interference?
- Are these task-specific updates additive — i.e., can we compose a multitask model post-hoc?

- What theoretical principles govern the geometry of continual updates in weight space?

Research Contributions

This work makes the following contributions:

1. We develop a unified framework to apply and evaluate LoRA-based continual learning in LLMs [4].
2. We adapt Orthogonal LoRA (O-LoRA) [11] to constrain task updates to orthogonal subspaces and analyze its performance.
3. We introduce Task Arithmetic [1] to enable post-hoc merging of task-specific updates and evaluate its generalization.
4. We perform comprehensive experiments on benchmark NLP tasks, showing reduced forgetting and improved scalability.

Catastrophic Forgetting in Continual Learning

Continual learning involves training a model on a sequence of tasks $\{T_1, T_2, \dots, T_n\}$ while maintaining performance on prior tasks. A central challenge in this setup is *catastrophic forgetting* — the tendency of a model to lose previously acquired knowledge when learning a new task [7].

This is illustrated in Figure 1.1, where training on a new task causes a sharp drop in accuracy on a previously learned one.

To address this, we explore *parameter-efficient fine-tuning* (PEFT) strategies—particularly **Low-Rank Adaptation (LoRA)** [4]—as a promising direction for continual learning. LoRA introduces small low-rank updates that adapt the model to new tasks while preserving the pre-trained weights.




Figure 1.1: Illustration of catastrophic forgetting in continual learning: performance on Task A drops sharply after training on Task B.

Building upon LoRA, we investigate two advanced methods:

- **Orthogonal LoRA (O-LoRA)** [11] — constrains each task’s updates to an orthogonal subspace to reduce interference.
- **Task Arithmetic** [1] — merges task-specific adapters using linear operations in weight space.

These methods enable scalable and cost-effective continual learning by reducing memory growth, training time, and catastrophic forgetting.

Background

In this chapter, we review foundational concepts that underpin continual learning in neural networks and large language models (LLMs). We begin by discussing the phenomenon of catastrophic forgetting and strategies to mitigate it, followed by a motivation for parameter-efficient fine-tuning in the context of scalable continual learning.

2.1 Continual Learning and Catastrophic Forgetting

Continual learning (CL), also referred to as lifelong or incremental learning, is the paradigm in which a model encounters a sequence of tasks $\{T_1, T_2, \dots, T_n\}$ and is expected to learn each new task without compromising performance on previously learned ones. This setting mimics real-world conditions where new data or objectives emerge over time.

A central challenge in CL is *catastrophic forgetting*, where gradient updates for new tasks override the information learned for previous tasks. In standard neural networks, fine-tuning the same parameters across multiple tasks leads to this destructive interference, as gradients for different tasks may point in conflicting directions. This problem was formally recognized by McCloskey and Cohen (1989) [7] and remains a core challenge in sequential learning.

An example is a language model trained to classify sports news (Task A) and later fine-tuned to classify political news (Task B). Without mechanisms for knowledge retention, performance on Task A typically drops after learning Task B. This phenomenon was

previously illustrated in Figure 1.1.

To mitigate forgetting, several classes of methods have been proposed:

- **Regularization-based methods** (e.g., Elastic Weight Consolidation or EWC [6]) constrain updates to parameters deemed important for previous tasks by adding penalty terms to the loss function.
- **Rehearsal or replay-based methods** store a subset of data from past tasks and replay them during training on new tasks, helping the model to remember older distributions.
- **Parameter isolation methods** allocate separate modules or subnetworks for each task, such as Progressive Networks [9].

Figure 2.1 illustrates each strategy. While these methods are effective in small-scale settings, they are difficult to scale to LLMs due to constraints on memory, compute, or implementation complexity.




Figure 2.1: Continual learning strategies: (a–b) regularization with Elastic Weight Consolidation (EWC), and (c) parameter-isolated architectural method (Progressive Networks).

2.2 Continual Learning in LLMs

As continual learning is applied to large-scale models, new challenges emerge. Unlike smaller networks, LLMs have billions of parameters, making them particularly sensitive to forgetting and interference. Simultaneously, their scale opens up new possibilities for subspace-based learning.

Continual learning for LLMs often follows task-incremental, class-incremental, or domain-incremental paradigms:

- **Task-incremental:** Task ID is known at test time.
- **Class-incremental:** New labels are introduced sequentially with no task ID.
- **Domain-incremental:** Same labels, but distributions shift over time.

Each of these setups presents unique scalability, memory, and generalization challenges. Figure 2.2 shows a categorization of continual learning paradigms for LLMs based on task visibility, label space, and input format.




Figure 2.2: Six categories of continual learning in LLMs. These differ based on task visibility, data format, and supervision signals.

2.3 Parameter-Efficient Fine-Tuning (PEFT)

Given the impracticality of full fine-tuning for each task, parameter-efficient fine-tuning (PEFT) techniques offer a promising alternative. PEFT methods adapt LLMs by modifying

only a small set of parameters or inserting lightweight modules while keeping the majority of the model frozen.

A prominent PEFT method is **Low-Rank Adaptation (LoRA)** [4], which injects low-rank updates into transformer layers. LoRA introduces a weight modification $\Delta W = AB$ to a frozen weight matrix W , where $A \in \mathbb{R}^{d \times r}$ and $B \in \mathbb{R}^{r \times k}$ and $r \ll \min(d, k)$.

This formulation injects only $2dr$ parameters per matrix — significantly reducing memory and compute costs while retaining strong adaptation capacity.




Figure 2.3: Common parameter-efficient fine-tuning (PEFT) methods including LoRA, DyLoRA, and DoRA. These methods reduce the number of trainable parameters per task.

Advantages of LoRA:

- Orders of magnitude fewer trainable parameters.
- No change to the model’s forward pass.
- LoRA adapters can be merged or composed additively.

2.4 Continual Learning with LoRA

LoRA’s structure makes it highly suitable for continual learning. However, naively applying LoRA in a sequential setting still leads to forgetting if the same adapter is reused across tasks. To address this, several recent methods have emerged:

- **Orthogonal LoRA (O-LoRA)** [11]: constrains each task’s adapter to be orthogonal to previous ones.

- **Task Arithmetic** [1]: merges task-specific adapters post-hoc using linear combinations.

These approaches preserve LoRA's parameter efficiency while introducing mechanisms to reduce task interference and improve knowledge retention.

The next chapter introduces the formulation and benefits of Orthogonal LoRA in detail.

Orthogonal Low-Rank Adaptation

In this chapter, we describe the Orthogonal Low-Rank Adaptation (O-LoRA) framework and our proposed modifications that improve its applicability for continual learning in large language models (LLMs). We begin by revisiting the formulation of LoRA and then explain how orthogonality constraints are introduced in O-LoRA to prevent catastrophic forgetting. Finally, we present empirical results showing its effectiveness.

3.1 Motivation and Challenges

Orthogonal Gradient Descent (OGD) [2] is a continual learning method that mitigates forgetting by constraining new gradients to be orthogonal to gradients from previous tasks. The key idea is to project the gradient of the current task onto the subspace orthogonal to those of previous tasks:

$$g_t^\perp = g_t - \sum_{i < t} \frac{g_t^\top g_i}{\|g_i\|^2} g_i$$

This ensures that new task updates do not interfere destructively with prior knowledge.

Orthogonal LoRA (O-LoRA) [11] extends this principle to the parameter-efficient setting by applying orthogonality to low-rank updates rather than full gradients. Instead of learning one shared adapter, each task is allocated a distinct low-rank update $\Delta W_t = A_t B_t$, constrained to avoid overlapping with updates from earlier tasks.

The goal is to enforce orthogonality across the task-specific update subspaces:

$$\text{span}(A_i) \perp \text{span}(A_t) \quad \text{for all } i < t$$

3.2 Orthogonal LoRA (O-LoRA)

Let $\Delta W_t = A_t B_t$ represent the low-rank update for task t . Then for two tasks i and t , orthogonality is expressed as:

$$\langle \Delta W_i, \Delta W_t \rangle = \text{Tr}[(A_i B_i)^\top (A_t B_t)] = 0$$

To implement this efficiently, O-LoRA replaces costly gradient projection (used in OGD) with a regularization-based approach during training. This encourages update directions to remain orthogonal by penalizing their inner products.

Orthogonality Regularization

Orthogonality loss term is defined as:

$$\mathcal{L}_\perp = \lambda \sum_{i < t} \|A_i^\top A_t\|_F^2$$

The total training objective becomes:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{task}} + \mathcal{L}_\perp$$

This ensures that the new task’s adapter update lies in a subspace orthogonal to previously learned directions.




Figure 3.1: Visualization of Orthogonal Gradient Descent (OGD): Gradients for new tasks are projected onto subspaces orthogonal to previously learned task gradients to reduce forgetting. O-LoRA adapts this principle using orthogonality between LoRA adapters.

3.3 O-LoRA Architecture

Figure 3.2 illustrates the O-LoRA setup. Each task introduces its own low-rank update to the base model. During training, regularization ensures that these updates remain orthogonal in their subspace structure, thus minimizing interference.

3.4 Empirical Results

We evaluate O-LoRA using the OPT-1.3B model across four NLP benchmarks under the sequential task learning setting [11]. These include:

- **DBpedia:** Topic classification
- **Amazon Reviews:** Sentiment analysis
- **Yahoo Answers:** Topic classification
- **AGNews:** News categorization




Figure 3.2: O-LoRA architecture: Each task-specific LoRA update is constrained to be orthogonal to previous tasks’ update directions, reducing interference and enabling sequential learning.

We compare O-LoRA to baseline methods: sequential LoRA (shared adapter), incremental LoRA (merged adapters), and the joint training upper bound.

Table 3.1: Performance of O-LoRA on benchmark datasets (OPT-1.3B)

Method	DBpedia	Amazon	Yahoo	AGNews	Avg. Acc.
Joint (Upper Bound)	98.67	73.37	74.43	89.57	84.01
Sequential LoRA	29.87	58.12	52.09	90.95	57.76
Incremental LoRA	31.03	58.58	56.12	89.71	58.86
O-LoRA ($\lambda = 0.1$)	45.50	47.66	55.57	90.25	59.75
O-LoRA ($\lambda = 0.5$)	68.02	38.60	56.93	90.23	63.44

Observations:

- O-LoRA significantly improves average accuracy over sequential baselines, especially at $\lambda = 0.5$.
- Even without replay or task-specific heads, O-LoRA effectively retains task performance across the sequence.
- Compared to naive LoRA, performance gain on DBpedia exceeds 35%, and average accuracy improves by 6%.

This confirms that orthogonalizing low-rank updates offers a scalable and effective route to continual learning.

In the next section, we explore extensions to O-LoRA that enforce orthogonality in the *effective update space* $\Delta W = AB$, rather than just in projection space.

3.5 Modified Orthogonal LoRA (M-LoRA)

While O-LoRA constrains the subspace spanned by the adapter input matrices A_t to be orthogonal, it does not guarantee orthogonality in the full update space $\Delta W = AB$. This limitation can result in residual interference, especially if the B matrices are unconstrained or aligned.

Key Insight: To eliminate this interference, we propose Modified O-LoRA (M-LoRA), which enforces orthogonality directly between the effective low-rank updates.

Instead of minimizing $\|A_i^\top A_t\|_F^2$, we now regularize the overlap between the full matrix products:

$$\mathcal{L}_\perp^{\text{M-LoRA}} = \lambda \sum_{i < t} \|(A_i B_i)^\top (A_t B_t)\|_F^2$$

The total loss becomes:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{task}} + \mathcal{L}_\perp^{\text{M-LoRA}}$$

This formulation penalizes overlapping task directions in the actual update space, yielding better generalization and robustness to forgetting.

Advantages over O-LoRA:

- Stronger constraint: applied in output space, not just projection subspace.
- Task separation is enforced directly on the model’s functional behavior.
- Reduced merging conflicts in subsequent continual learning steps.




Figure 3.3: M-LoRA extends O-LoRA by constraining the actual task-specific update matrices (AB) to lie in orthogonal subspaces.

3.6 Empirical Results

We evaluated both O-LoRA and M-LoRA on the OPT-1.3B model across four sequential NLP tasks. The table below summarizes the accuracy per task and average accuracy for each method.

Table 3.2: Performance of sequential adaptation methods on benchmark datasets (OPT-1.3B)

Method	DBpedia	Amazon	Yahoo	AGNews	Avg. Acc.
Sequential LoRA	29.87	58.12	52.09	90.95	57.76
O-LoRA ($\lambda = 0.5$)	68.02	38.60	56.93	90.23	63.44
Modified O-LoRA (M-LoRA)	72.86	48.11	58.68	89.84	67.37

Key Observations:

- M-LoRA outperforms O-LoRA and Sequential LoRA across all tasks, with improvements most prominent on DBpedia and Amazon.
- Average accuracy improves by nearly 4% over O-LoRA and by 10% over the naive sequential baseline.

- Cosine similarity between task vectors confirms stronger orthogonality under M-LoRA, validating our hypothesis.

Emergent Orthogonality in Task Vectors

Interestingly, even in the absence of explicit orthogonality regularization, task-specific LoRA updates often exhibit near-orthogonal behavior in high-dimensional weight space [5]. To quantify this, we compute the cosine similarity between update matrices across tasks:

$$\cos(\theta_{i,j}) = \frac{\langle A_i B_i, A_j B_j \rangle}{\|A_i B_i\| \cdot \|A_j B_j\|}$$

Empirical results show that $\cos(\theta_{i,j}) \approx 0$ for most (i, j) pairs, suggesting that the directions occupied by individual task vectors are already weakly correlated or nearly orthogonal.




Figure 3.4: Cosine similarity distribution among task vectors $A_i B_i$. The majority cluster around zero, indicating emergent orthogonality.

Theoretical Justification via Random Vectors

Let $X, Y \in \mathbb{R}^d$ be independent, isotropic random vectors with sub-Gaussian coordinates and variance 1. Then, for universal constants $c, C > 0$ and any $\varepsilon > 0$,

$$\Pr(|\cos(\theta)| \geq \varepsilon) \leq C \exp(-c\varepsilon^2 d)$$

This result implies that the probability of two random vectors being significantly non-orthogonal drops exponentially with the dimension d [10]. Therefore, in practice:

- Cosine similarity between two random directions is typically close to zero.
- The angle between high-dimensional vectors concentrates sharply around 90° .
- Any fixed set of $k \leq e^{\alpha d}$ such vectors are likely to be pairwise ε -orthogonal with high probability.

These theoretical properties provide a foundation for understanding why LoRA task vectors — which operate in high-dimensional weight space — often exhibit near-orthogonality even without enforced constraints. This supports the design rationale behind methods like O-LoRA and M-LoRA that formalize and strengthen this behavior through regularization.

Task Arithmetic for Model Merging

While Orthogonal LoRA focuses on sequential task adaptation with interference reduction, an alternate approach known as *Task Arithmetic* treats each task’s weight update as an independent vector in parameter space. The core idea is that these task-specific updates — when learned separately — can be algebraically combined to construct a multi-task model post hoc, sometimes without needing to retrain [5, 1].

4.1 The Concept of Task Vectors

Let θ_0 denote the original pre-trained model weights, and θ_t^* denote the fine-tuned model weights after training on task t . The difference vector:

$$\tau_t = \theta_t^* - \theta_0$$

is referred to as the **task vector**. It represents how task t shifts the model in weight space [5].

In the context of LoRA [4], each task’s update is a low-rank approximation:

$$\Delta W_t = A_t B_t \quad \Rightarrow \quad \tau_t \approx A_t B_t$$

After learning tasks independently, these updates can be merged:

$$\theta_{\text{merged}} = \theta_0 + \sum_{t=1}^n \alpha_t \tau_t$$

where α_t are interpolation weights (default $\alpha_t = 1$).




Figure 4.1: Task Arithmetic concept: Individual task-specific LoRA updates are merged to create a single model that performs well on all tasks.

Motivation:

- Avoids sequential training and forgetting.
- Parallel training of tasks possible.
- Enables on-demand multi-task composition.

4.2 Theoretical Insights on Task Disentanglement

Empirical success of task arithmetic is often attributed to an emergent property of pre-trained models — **weight disentanglement**. Ilharco et al. [5] formalize this through the following expression:

$$f\left(x; \theta_0 + \sum_t \alpha_t \tau_t\right) = \sum_t g_t(x; \alpha_t \tau_t) + g_0(x)$$

where each g_t influences a distinct region of the input space.

In practice, we find that task vectors behave approximately orthogonally in high-dimensional weight space, especially when trained on disjoint tasks [10]. This forms the basis for effective merging with minimal conflict.




Figure 4.2: Disentangled task vector behavior: Separate task-specific directions influence disjoint input regions, enabling effective merging via Task Arithmetic.

4.3 Orthogonality of Task Vectors

We define the merging interference for each task i using the *merging gap* G_i :

$$G_i = L_i \left(\theta + \sum_j \lambda_j \tau_j \right) - L_i(\theta + \lambda_i \tau_i)$$

where L_i is the loss on task i , and λ_j are interpolation scalars. A first-order Taylor expansion gives [5]:

$$G_i \approx k_i \sum_{j \neq i} \lambda_j \langle \tau_i, \tau_j \rangle$$

where $k_i < 0$. Thus, if task vectors are orthogonal, i.e., $\langle \tau_i, \tau_j \rangle = 0$ for $i \neq j$, we have $G_i \approx 0$. This supports the need for orthogonal task representations during training.

4.4 Orthogonality Regularization for Task Arithmetic

To ensure successful merging, task vectors should ideally lie in orthogonal subspaces [11].

This motivates explicit orthogonality regularization:

$$\mathcal{L}_{\perp}^{\text{TA}} = \lambda \sum_{i < t} \left(\frac{\tau_i^{\top} \tau_t}{\|\tau_i\| \|\tau_t\|} \right)^2$$

This term is added to the training loss of each task, encouraging updates that are geometrically independent of earlier tasks.

Merged Model Construction

After training all tasks with regularization, we form a merged model as:

$$\theta_{\text{merged}} = \theta_0 + \sum_{t=1}^T \alpha_t \tau_t$$

Benefits of Orthogonal Task Arithmetic:

- Prevents destructive interference.
- Boosts post-merge accuracy.
- Compatible with memory-efficient deployment.

4.5 Empirical Results

We evaluate task arithmetic using OPT-1.3B on four NLP benchmarks: DBpedia, Amazon Reviews, Yahoo Answers, and AGNews. Models were trained independently on each task using LoRA adapters, with and without orthogonal regularization. Task vectors were




Figure 4.3: Comparison of cosine similarity heatmaps of task vectors with and without orthogonality regularization. Regularization (a) reduces overlap between task vectors, improving merge stability; without it (b), interference is more likely [11].

merged using weighted addition:

$$\theta^* = \theta_0 + \alpha \sum_{t=1}^T A_t B_t$$

Table 4.1: Performance of Task Arithmetic (OPT-1.3B) with varying interpolation weight α

Method	DBpedia	Amazon	Yahoo	AGNews	Avg. Acc.
Single Task (LoRA)	98.66	73.06	72.47	88.09	83.07
Orth. Reg. STL	98.66	59.07	73.21	88.33	79.82
TA ($\alpha = 0.25$)	64.53	39.42	50.02	62.14	54.03
TA + Orth. Reg. ($\alpha = 0.25$)	70.29	38.68	56.10	64.34	57.35
TA ($\alpha = 0.3$)	71.15	37.72	54.94	57.51	55.33
TA ($\alpha = 0.4$)	77.30	28.53	57.88	55.26	54.74
TA ($\alpha = 0.5$)	74.51	5.53	55.43	52.76	47.06

Observations:

- Orthogonal regularization improves average accuracy by over 3% compared to unregularized merging (Table 4.1).
- As seen in Figure 4.3, cosine similarity between task vectors is significantly reduced when orthogonality regularization is applied [11].

- Without regularization, overlapping update directions lead to performance drops.
- Merging performance is also sensitive to the interpolation weight α . Small values (e.g., 0.25–0.3) yield optimal trade-offs; beyond $\alpha = 0.4$, instability and accuracy drop emerge.
- Task arithmetic approaches joint training performance without retraining or access to all tasks simultaneously [5, 1].

Experimental Setup and Results

This chapter describes the experimental protocols used to evaluate our proposed approaches for continual learning in LLMs. We provide implementation details, dataset information, training configuration, and report performance across multiple benchmarks. The evaluation focuses on three main goals: knowledge retention (resistance to forgetting), effective acquisition of new tasks, and parameter efficiency.

5.1 Model and Framework

We use the OPT-1.3B model, a decoder-only transformer architecture pretrained on large-scale internet text using causal language modeling [?]. OPT is similar in structure to GPT and is open-sourced by Meta AI.

- **Model size:** 1.3 billion parameters
- **Tokenizer:** GPT2Tokenizer
- **Adapters:** LoRA adapters inserted at attention projection layers (query and value) [4]
- **Rank:** $r = 8$ (unless otherwise specified)

Training is conducted using:

- **DeepSpeed** and **Accelerate** for multi-GPU distributed training on NYU HPC [?, ?]

- **Mixed precision (fp16)** for memory efficiency
- **Batch size:** 32
- **Learning rate:** $1e-4$ with cosine decay

5.2 Datasets and Tasks

We experiment with four popular NLP classification datasets, each treated as a separate task in a sequential continual learning setup:

- **AGNews:** Topic classification (4 categories) [?]
- **Amazon Reviews:** Sentiment classification (5 levels) [?]
- **Yahoo Answers:** Topic classification (10 categories) [?]
- **DBpedia:** Ontology classification (14 categories) [?]

Each task is trained sequentially and evaluated immediately after training, as well as after learning all tasks. We follow a task-incremental setting where task boundaries are known but shared output heads are used.

5.3 Evaluation Metrics

To evaluate continual learning performance, we report:

- **Task Accuracy:** Final accuracy on each task after completing the full training sequence.
- **Average Accuracy:** Mean accuracy across all tasks after final training.
- **Forgetting:** Drop in accuracy from maximum to final value for each task.

- **Memory Usage:** Total additional parameters used for adaptation (measured in millions).

5.4 Method Comparison

We compare the following methods:

- **Sequential Fine-Tuning:** Updates the full model per task [7].
- **Sequential LoRA:** Reuses the same LoRA adapter across tasks [4].
- **O-LoRA:** Applies orthogonality constraints to LoRA adapters (A matrices) [11].
- **Modified O-LoRA (M-LoRA):** Enforces orthogonality in the effective update space $\Delta W = AB$.
- **Task Arithmetic (TA):** Merges independently trained LoRA updates using weighted addition [5].
- **TA + Orth. Regularization:** Adds cosine-based regularization during training to ensure orthogonality of task vectors [11].
- **Joint Training (Upper Bound):** Trained on all tasks simultaneously.

5.5 Performance Summary

Key Observations

Based on the accuracy scores summarized in Table 5.1, we highlight the following insights:

- **Joint Training** unsurprisingly achieves the highest average accuracy of 84.01%, serving as the upper bound since it benefits from full multitask exposure and shared optimization.

Table 5.1: Final accuracy (%) per task and average across methods (OPT-1.3B)

Method	DBpedia	Amazon	Yahoo	AGNews	Avg. Acc.
Joint Training (Upper Bound)	98.67	73.37	74.43	89.57	84.01
Sequential Fine-Tune	27.15	58.72	52.43	91.32	57.91
Sequential LoRA	29.87	58.12	52.09	90.95	57.76
O-LoRA ($\lambda = 0.5$)	68.02	38.60	56.93	90.23	63.44
Modified O-LoRA	72.86	48.11	58.68	89.84	67.37
TA ($\alpha = 0.25$)	64.53	39.42	50.02	62.14	54.03
TA + Orth. Reg. ($\alpha = 0.25$)	70.29	38.68	56.10	64.34	57.35

- **Sequential Fine-Tuning** and **Sequential LoRA** suffer significantly from catastrophic forgetting, achieving only around 57–58% average accuracy despite high performance on the last task (AGNews).
- **O-LoRA** ($\lambda = 0.5$) substantially improves upon sequential LoRA by enforcing subspace separation via orthogonality constraints [11], leading to a 5.7% absolute gain in average accuracy.
- **Modified O-LoRA (M-LoRA)** further boosts performance by applying orthogonality in the effective update space ($\Delta W = AB$), achieving an average accuracy of 67.37% — the highest among continual methods.
- **Task Arithmetic (TA, $\alpha = 0.25$)** without regularization performs poorly (54.03%), likely due to task vector interference during merging [5].
- **TA + Orthogonal Regularization** mitigates this issue by enforcing cosine separation between task vectors during training [11]. This yields a 3.3% improvement in average accuracy, raising it to 57.35%. This supports our hypothesis that orthogonality enables more reliable post-hoc merging.
- **Comparison with O-LoRA:** While TA + Orth. Reg. performs better than naive TA, it still falls short of O-LoRA and M-LoRA, suggesting that sequential subspace-aware adaptation remains more robust than pure post-hoc merging under current conditions.

Discussion

In this chapter, we interpret the empirical results and analyze the practical implications of our proposed low-rank continual learning techniques. We reflect on the theoretical motivations, method trade-offs, and provide guidelines for real-world deployment of these approaches.

6.1 Stability vs. Plasticity Trade-Off

A central theme in continual learning is balancing:

- **Stability** — the ability to retain knowledge from previous tasks.
- **Plasticity** — the flexibility to acquire new task knowledge.

Our experiments highlight this trade-off:

- **Sequential Fine-Tuning** achieves high plasticity but suffers severely from catastrophic forgetting [7].
- **Sequential LoRA** improves memory efficiency [4] but does not adequately prevent forgetting.
- **O-LoRA and Modified O-LoRA** introduce orthogonal constraints [11] to preserve task-specific updates and strike a stronger balance.

- **Task Arithmetic with Orthogonal Regularization** enables post-hoc multi-task composition with significantly reduced interference, achieving performance competitive with multitask joint training [5].

6.2 Orthogonality in Weight Space

One of the most impactful theoretical insights of this work is the emergent **approximate orthogonality** of task-specific low-rank update directions in high-dimensional parameter space.

- Even without explicit constraints, task vectors in pre-trained LLMs exhibit low cosine similarity — a result supported both empirically and by theory (via random vector concentration properties) [10].
- **Modified O-LoRA** enhances this structure by enforcing orthogonality in the effective update space $\Delta W = AB$, resulting in better separation and improved retention across tasks [11].
- In **Task Arithmetic**, orthogonality regularization significantly improves post-merge accuracy and reduces cosine similarity between task vectors, as shown in Figure ??.

6.3 Effectiveness of Task Arithmetic

Task Arithmetic demonstrates that task-specific LoRA updates can be merged algebraically:

$$\theta_{\text{merged}} = \theta_0 + \alpha \sum_{t=1}^T A_t B_t$$

- When trained independently, task vectors τ_t can be combined without retraining, yielding a multitask model [5, 1].

- However, without regularization, overlapping task directions lead to destructive interference and degraded accuracy.
- Adding orthogonality constraints ($\mathcal{L}_\perp$) during training leads to more disentangled updates and improved performance [11].

The average accuracy improves from 54.03% (**TA** $\alpha = 0.25$ without regularization) to 57.35% (**TA + Orth. Reg.**), supporting the utility of this simple yet powerful merging strategy.

6.4 Comparative Observations

From the combined evaluation table, we observe:

- **Modified O-LoRA** achieves the best balance among continual learning methods, outperforming naive sequential LoRA and classic O-LoRA with an average accuracy of 67.37%.
- **Task Arithmetic with Orthogonal Regularization** achieves 57.35% average accuracy, even without sequential training — underscoring the promise of parallel, memory-efficient training [5].
- **Joint Training**, as expected, performs best (84.01%) but requires full access to all data and is not a feasible continual learning method.

6.5 Guidelines for Method Selection

- Use **O-LoRA** when tasks must be learned sequentially and replay or task storage is not allowed.
- Use **Modified O-LoRA** when orthogonality in the output space is desired and full task-specific updates should be disentangled.

- Use **Task Arithmetic + Orthogonal Regularization** for memory-efficient post-hoc multi-task composition from independently trained tasks.

6.6 Limitations and Future Work

While promising, our methods have limitations:

- Orthogonality regularization increases training time, especially as the number of tasks grows [8].
- Task vectors occupy memory equal to the LoRA updates (AB), which may become significant for 100+ tasks.
- We have not explored continual domain adaptation or few-shot settings.

Future extensions may include:

- **SD-LoRA**: Scaling task arithmetic using normalized task vectors and learnable interpolation weights [1].
- **Unsupervised task inference**: Determining task identity dynamically at inference time.
- **Theoretical analysis**: Deeper understanding of geometry in LoRA subspaces and interference mechanisms.

Conclusion and Future Work

7.1 Conclusion

This thesis explored parameter-efficient continual learning strategies for large language models (LLMs), focusing on low-rank adaptation techniques and task arithmetic for model merging. Our key contributions are:

- **Orthogonal Low-Rank Adaptation (O-LoRA):** We introduced O-LoRA, which applies orthogonality constraints to LoRA adapters, effectively mitigating catastrophic forgetting in sequential task learning scenarios [11].
- **Modified O-LoRA (M-LoRA):** Building upon O-LoRA, M-LoRA enforces orthogonality in the effective update space ($\Delta W = AB$), leading to improved task separation and higher average accuracy across tasks.
- **Task Arithmetic with Orthogonal Regularization:** We demonstrated that applying orthogonality regularization during independent task training enhances the effectiveness of task arithmetic in merging models, reducing interference and improving performance [5].
- **Comprehensive Empirical Evaluation:** Through extensive experiments on NLP benchmarks such as DBpedia, Amazon Reviews, Yahoo Answers, and AGNews, we validated the efficacy of our proposed methods. Notably, M-LoRA achieved an average accuracy of 67.37%, outperforming other continual learning approaches.

7.2 Future Work

While our methods show promise, several avenues remain for further exploration:

- **Dynamic Task Weighting in Task Arithmetic:** Investigate adaptive schemes for determining interpolation weights (α_t) during model merging to optimize performance across tasks [1].
- **Scalability to Numerous Tasks:** Assess the scalability of O-LoRA and task arithmetic methods when applied to a larger number of tasks, ensuring efficiency and performance are maintained [8].
- **Integration with Other PEFT Techniques:** Explore the combination of our orthogonality-based methods with other parameter-efficient fine-tuning techniques, such as adapter fusion [3] or prompt tuning, to enhance adaptability.
- **Application to Multimodal Models:** Extend the application of orthogonal low-rank adaptation and task arithmetic to multimodal models, addressing the challenges of continual learning in diverse data modalities.
- **Theoretical Analysis of Orthogonality Constraints:** Conduct a deeper theoretical investigation into the role of orthogonality in mitigating interference and its impact on the geometry of the parameter space [10].

By addressing these areas, we aim to further advance the field of continual learning in large language models, making them more adaptable and efficient for real-world applications.

Bibliography

- [1] Abhishek Chitale, Damjan Kici, Omar Shaikh, Jiayuan Park, Mor Geva, Edward Beeching, Matt Gardner, and Douwe Kiela. Task arithmetic: Emergent properties in the weight space of language models. In *arXiv preprint arXiv:2310.02288*, 2023.
- [2] Mehrdad Farajtabar, Navid Azizan, Alex Mott, and Ang Li. Orthogonal gradient descent for continual learning. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2020.
- [3] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea Gesmundo, Mohammad Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for nlp. In *International Conference on Machine Learning (ICML)*, pages 2790–2799, 2019.
- [4] Edward Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.
- [5] Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Suchin Gururangan, Ludwig Schmidt, Hannaneh Hajishirzi, and Ali Farhadi. Editing models with task arithmetic. In *arXiv preprint arXiv:2212.04089*, 2023.
- [6] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Krzysztof Milan, Raia Hadsell, Nicolas Heess, Koray

- Kavukcuoglu, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- [7] Michael McCloskey and Neal J. Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24:109–165, 1989.
- [8] Yifan Qiao, Zexue Jiang, Zhiwen He, Yongsheng Ge, and Weizhu Liu. A closer look at low-rank adaptation for continual learning: Stability-plasticity trade-off and rank growth. *arXiv preprint arXiv:2402.16176*, 2024.
- [9] Andrei A Rusu, Neil C Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. *arXiv preprint arXiv:1606.04671*, 2016.
- [10] Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.
- [11] Xiaojie Wang, Wenhui Li, Shiyang Wang, Pengfei Liu, and Amanpreet Singh. Orthogonal low-rank adaptation for continual learning in language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.